

Processes

Chapter 3

Processes


- Process Concept
- Process Scheduling
- Operations on Processes
- Inter-process Communication
- Communication in Client-Server Systems

Process Concept

- **Process** – a program in execution; process execution must progress in sequential fashion
- Program is *passive* entity stored on disk (**executable file**), process is *active*
 - Program becomes process when executable file loaded into memory
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider multiple users executing the same program

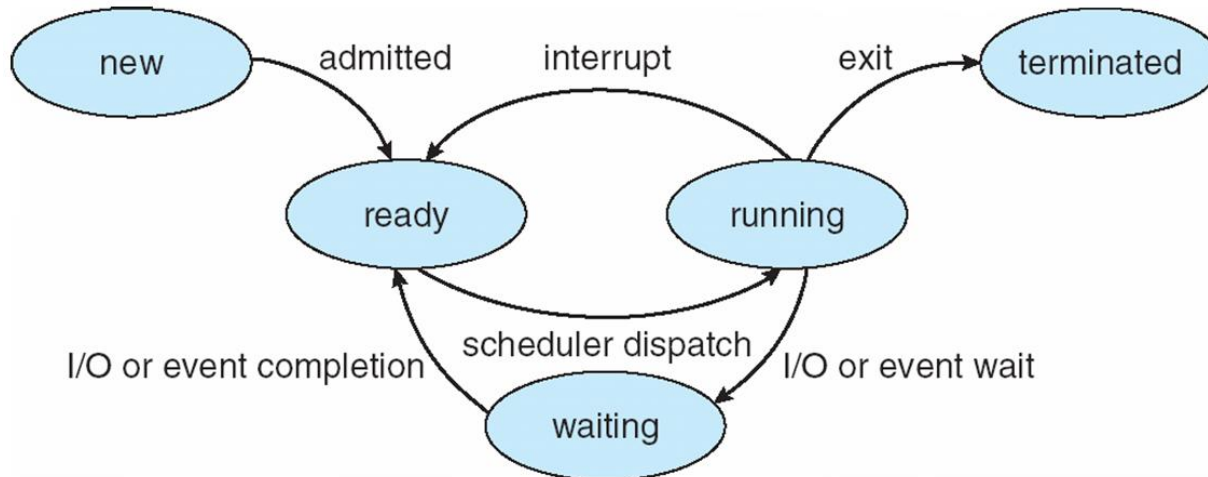
Process in Memory

- Multiple parts of process memory space
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time

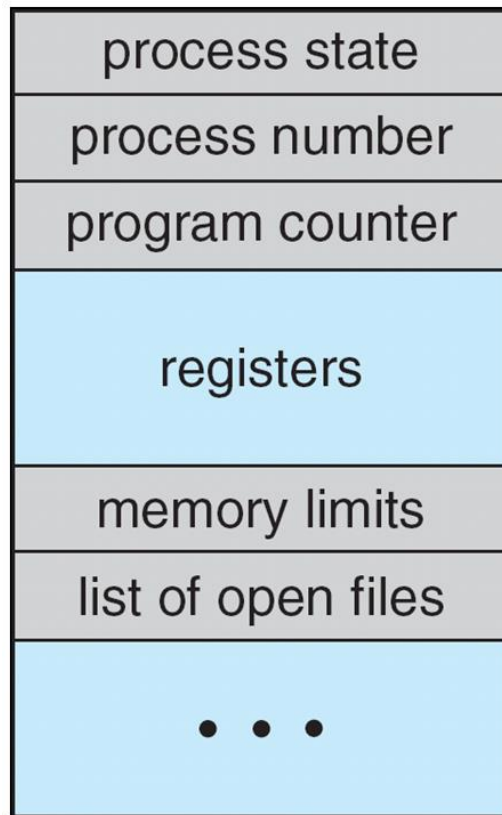


Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution



Process Control Block (PCB)



- Information associated with each process
 - Process state – running, waiting, etc
 - Program counter – location of instruction to next execute
 - CPU registers – contents of all process-centric registers
 - CPU scheduling information- priorities, scheduling queue pointers
 - Memory-management information – memory allocated to the process
 - Accounting information – CPU used, clock time elapsed since start, time limits
 - I/O status information – I/O devices allocated to process, list of open files
- Threads

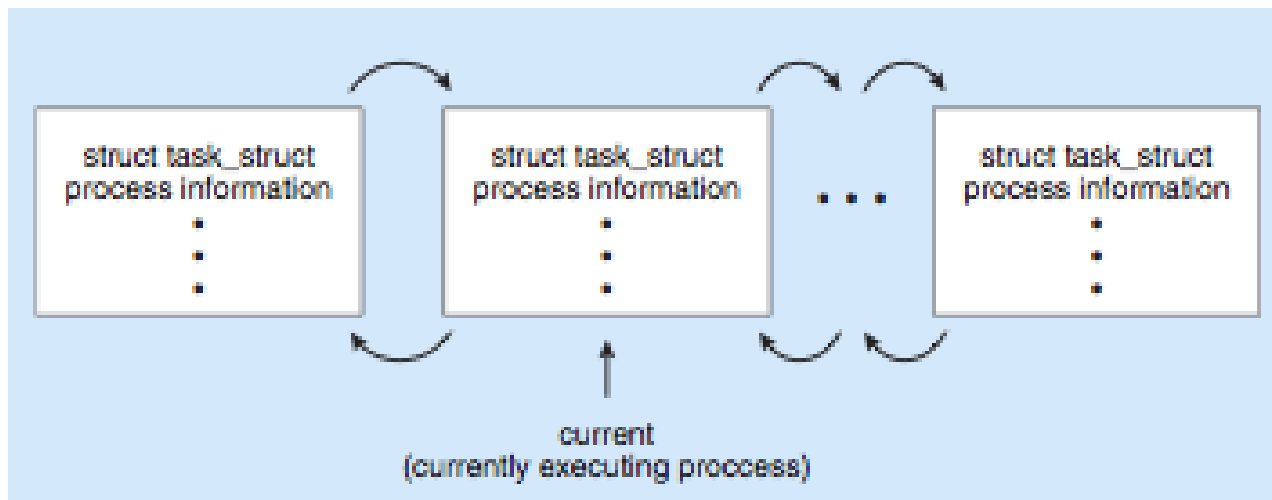
Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB
- More in next chapter

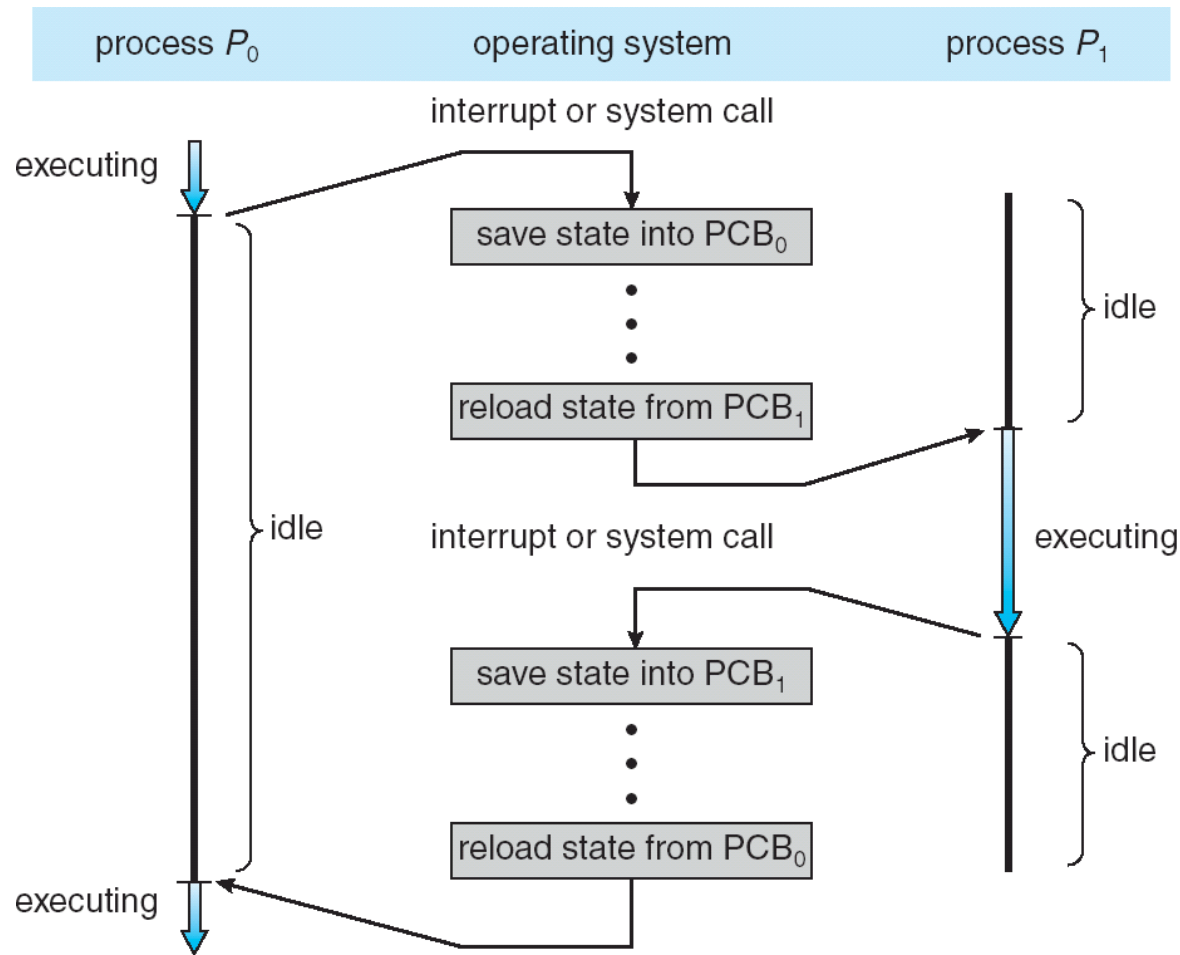
Process Representation in Linux

- Represented by the C structure `task_struct`

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



CPU Switch From Process to Process



Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process represented in the PCB
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> longer the context switch
- Time dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once

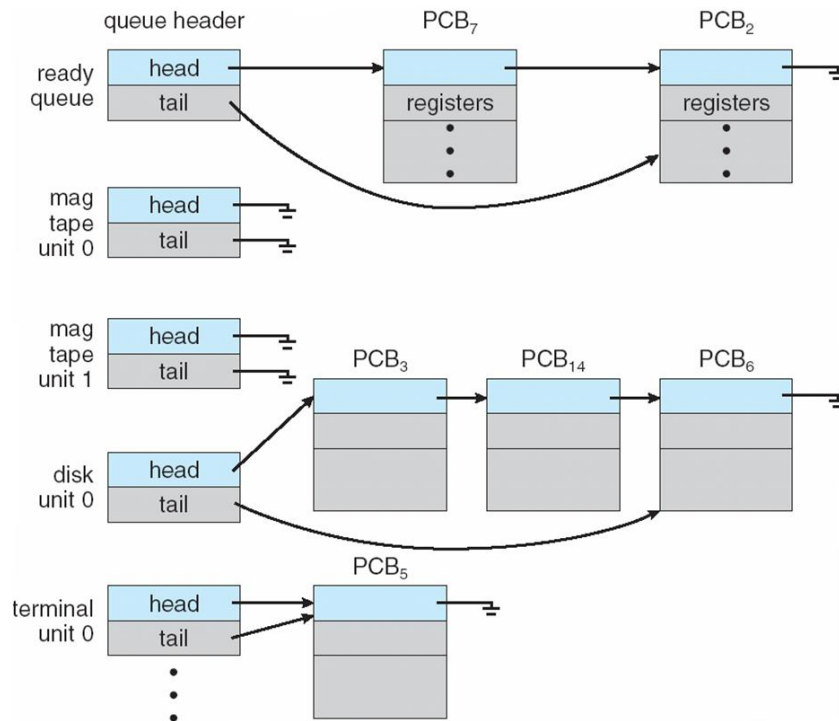
Process Context

- Which of following is NOT part of process context

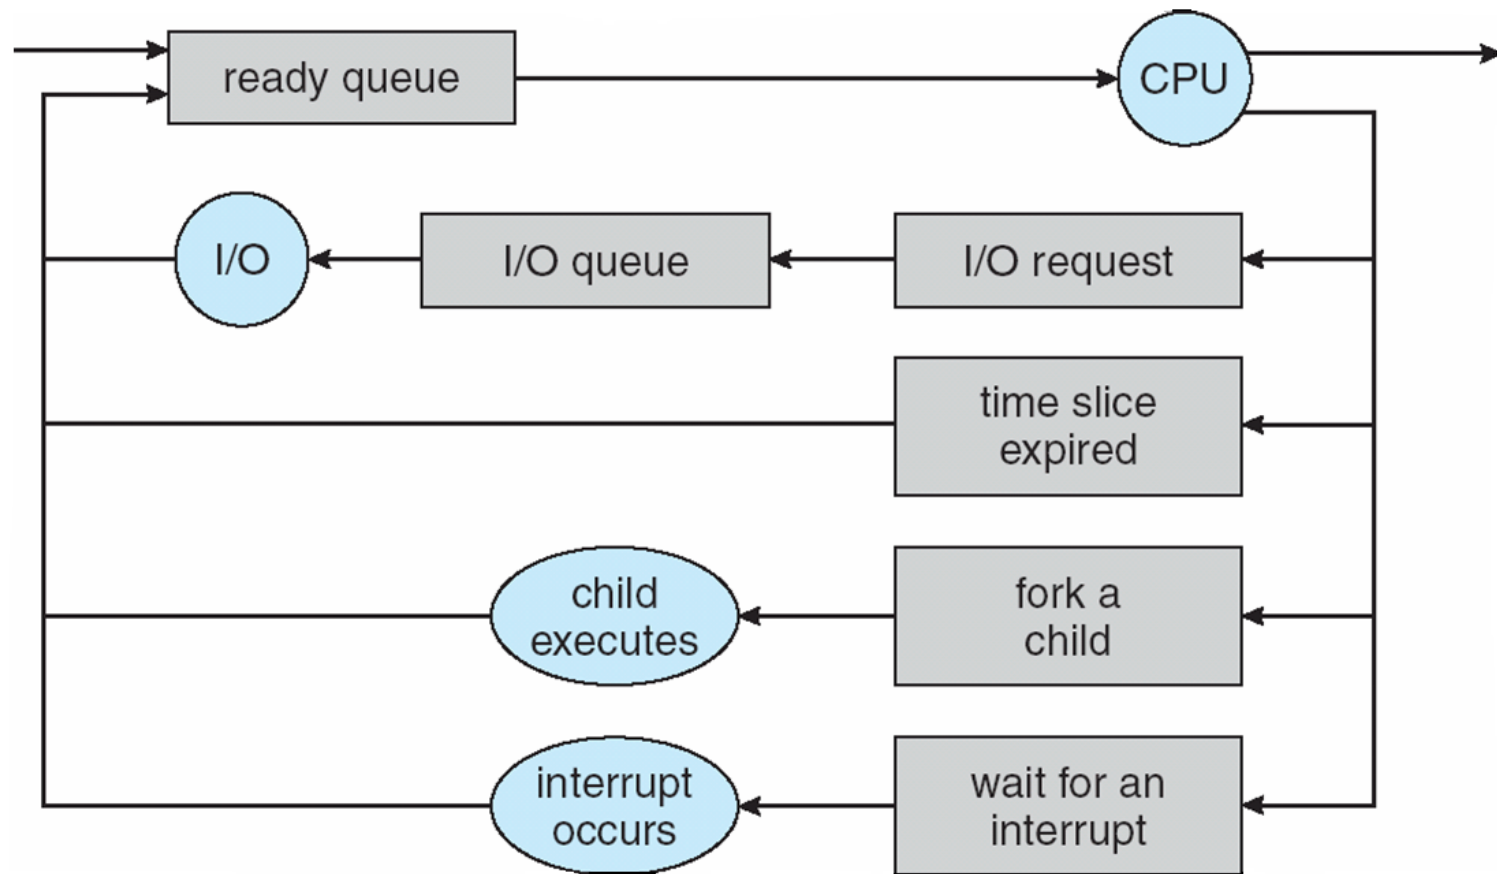
A	Program Counter
B	Register Values
C	Data section
D	Memory-management information
E	All of above are in process context

Process Scheduling Queues

- **Job queue** (process table) – set of all processes in the system
- **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- **Device queues** – set of processes waiting for an I/O device
- Processes migrate among the various queues



Representation of Process Scheduling



Schedulers

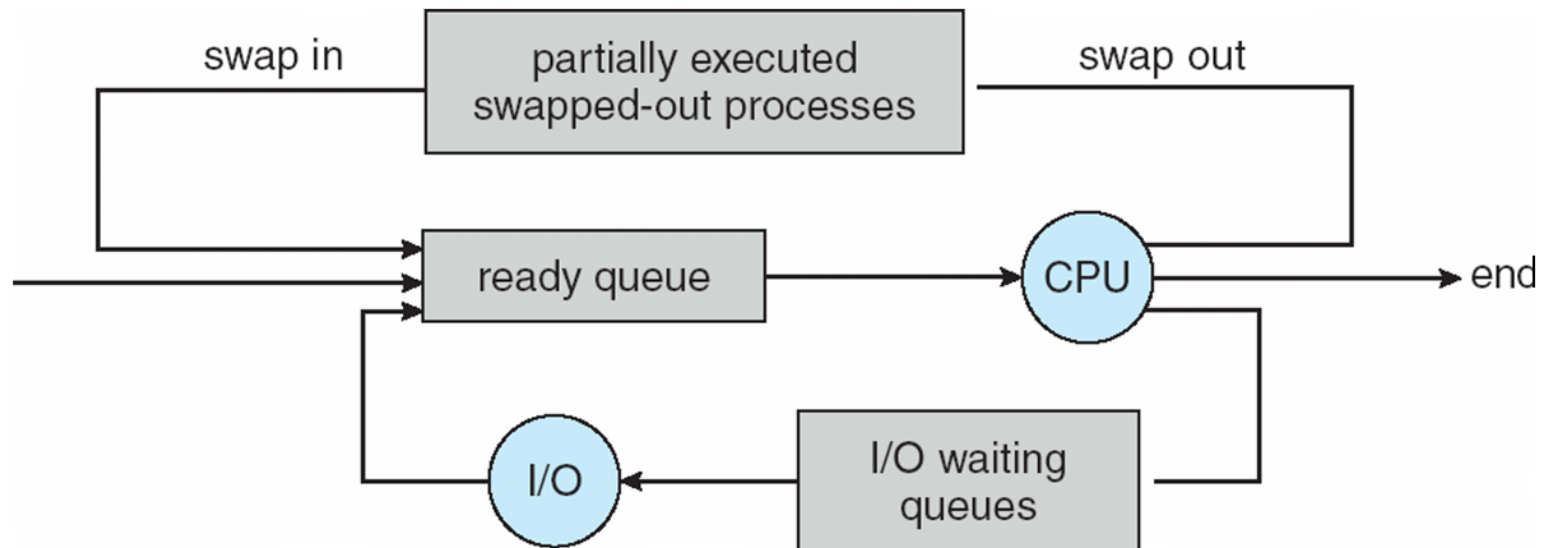
- OS uses schedulers to select processes from various queues.
- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the system
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU

Schedulers (Cont.)

- Short-term scheduler is invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*
- Processes can be described as either:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts
 - Long-term scheduler strives for good process mix

Addition of Medium Term Scheduling

- **Medium-term scheduler** can be added if degree of multiple programming needs to decrease
 - Remove process from memory, store on disk, bring back in from disk to continue execution: **swapping**



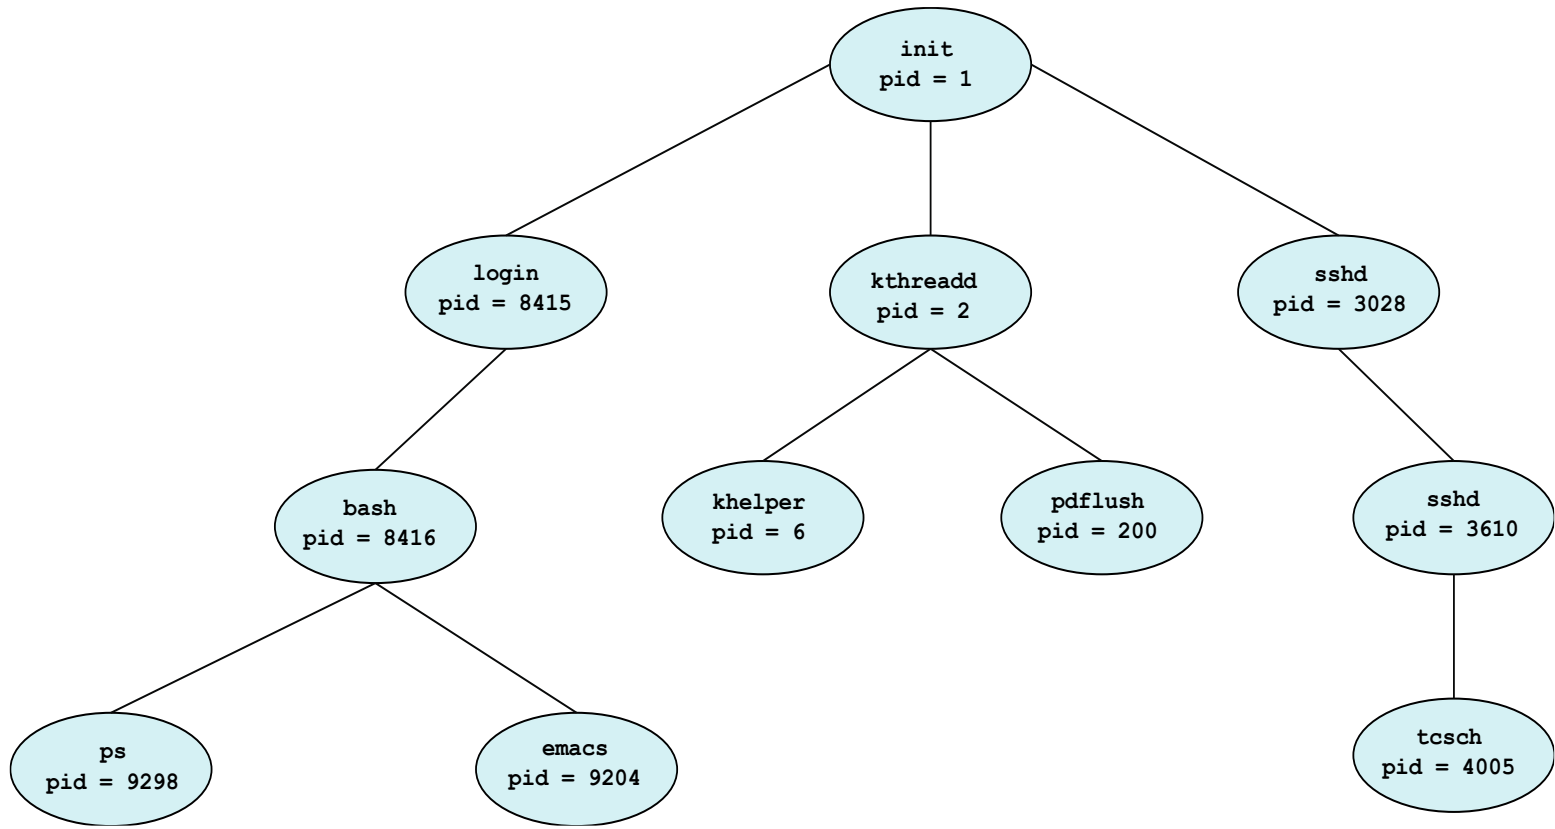
Multitasking in Mobile Systems

- Some systems / early systems allow only one process to run, others suspended
- Due to screen real estate, user interface limits iOS provides for a
 - Single **foreground** process- controlled via user interface
 - Multiple **background** processes– in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use

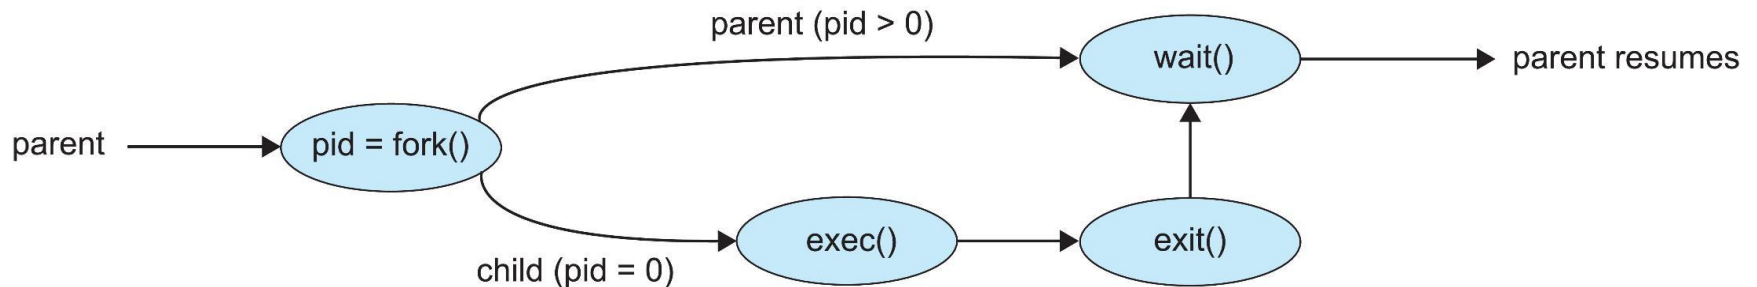
Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

A Tree of Processes in Linux



Process Creation (Cont.)



- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - **fork()** system call creates new process
 - **exec()** system call used after a **fork()** to replace the process' memory space with a new program

Forking Example p1.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
    } else {
        // parent goes down this path (original process)
        printf("hello, I am parent of %d (pid:%d)\n", rc, (int) getpid());
    }
    return 0;
}
```

Forking Example p2.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[])
{
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {
        // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        sleep(1);
    } else {
        // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n", rc, wc, (int) getpid());
    }
    return 0;
}
```

Forking Example p3.c

```
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    printf("hello world (pid:%d)\n", (int) getpid());
    int rc = fork();
    if (rc < 0) {
        // fork failed; exit
        fprintf(stderr, "fork failed\n");
        exit(1);
    } else if (rc == 0) {    // child (new process)
        printf("hello, I am child (pid:%d)\n", (int) getpid());
        char *myargs[3];
        myargs[0] = "wc"; // program: "wc" (word count)
        myargs[1] = "p3.c"; // argument: file to count
        myargs[2] = NULL; // marks end of array
        execvp(myargs[0], myargs); // runs word count
        printf("this shouldn't print out\n");
    } else {    // parent goes down this path (original process)
        int wc = wait(NULL);
        printf("hello, I am parent of %d (wc:%d) (pid:%d)\n",
               rc, wc, (int) getpid());
    }
    return 0;
}
```

Practice Question

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int value = 5;

int main()
{
    pid_t pid;

    pid = fork();

    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value = %d",value); /* LINE A */
        return 0;
    }
}
```

What output will be at Line A

A	PARENT: value = 0
B	PARENT: value = 5
C	PARENT: value = 15
D	PARENT: value = 20
E	None of Above

Figure 3.30 What output will be at Line A?

Practice Question

```
#include <stdio.h>
#include <unistd.h>

int main()
{
    /* fork a child process */
    fork();

    /* fork another child process */
    fork();

    /* and fork another */
    fork();

    return 0;
}
```

How many processes are created, including the parent process?

A	3
B	4
C	8
D	12
E	None of Above

Figure 3.31 How many processes are created?

Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**), e.g. **exit(0)**
 - Process' resources are deallocated by operating system
 - Returns status data from child to parent (via **wait()**)
- The parent may wait for a child to terminate using the **wait()** system call.

```
int cpid_done, status;
```

```
cpid_done = wait(&status);
```

- The caller sleeps until any child terminates.
- cpid_done gets the pid of the child that terminated.

Process Termination

- Parent may terminate execution of children processes
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
- If parent is exiting
 - Some operating system do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*
- If no parent waiting (did not invoke **wait()**), a terminated process becomes a **zombie**
- If parent terminated without invoking **wait**, a child process is an **orphan**

Basic Shell code

Pseudo-code a basic shell (not real C++ code):

```
while (1) {  
    parse_command_line; //get command, args, redirect,  
    etc.  
    if(cmd == exit) exit();  
    p = fork( );  
    if (p == 0) {  
        execvp (cmd, args)  
    } else {  
        if (command doesn't end with &)  
            wait( );  
    }  
}
```

Multi-process Architecture – Chrome Browser

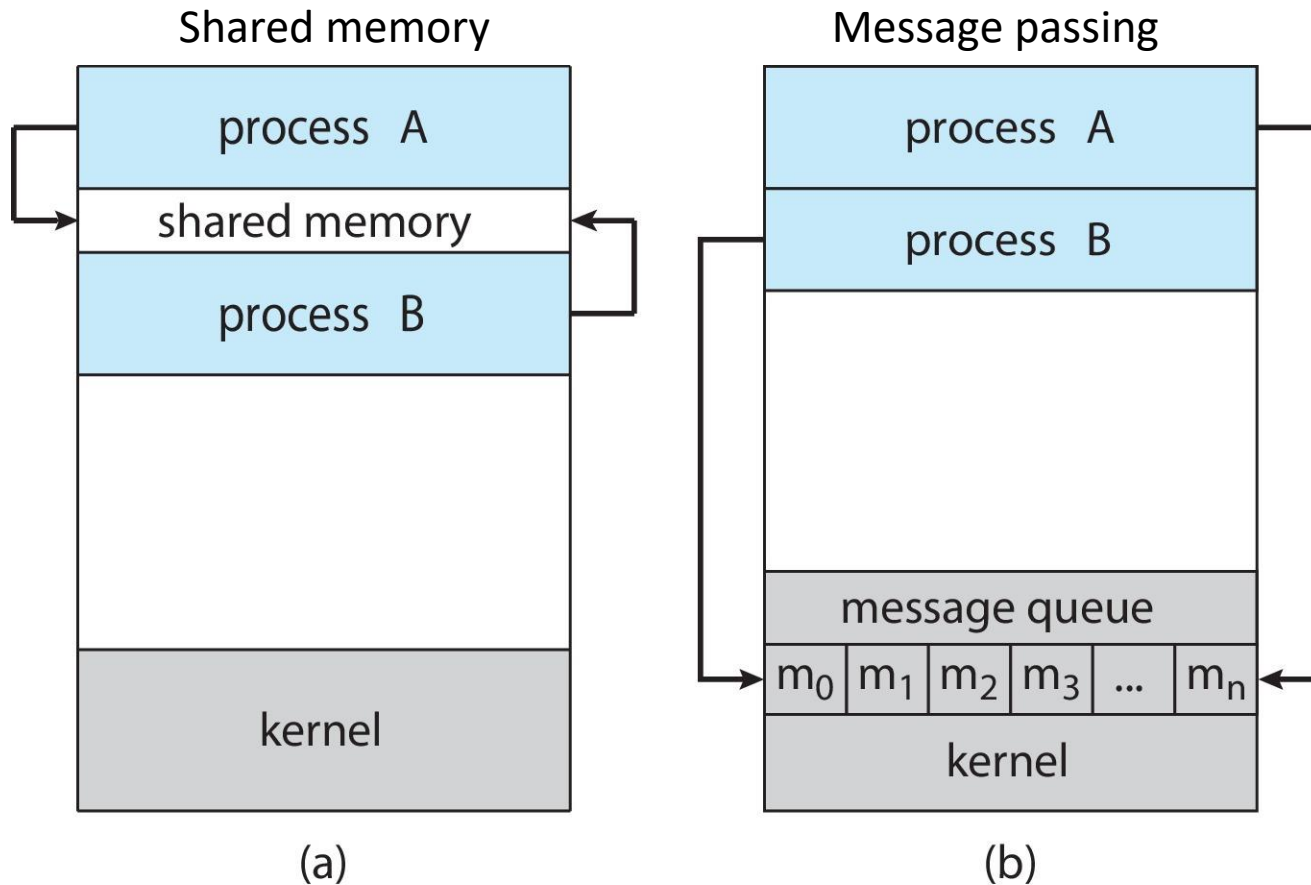
- Many web browsers ran as single process (some still do)
 - If one web site causes trouble, entire browser can hang or crash
- Google Chrome Browser is multi-process with 3 different types of processes:
 - **Browser** process manages user interface, disk and network I/O. A new browser process is created when Chrome is started
 - **Renderer** process renders web pages, deals with HTML, Javascript. A new renderer created for each website opened
 - Runs in **sandbox** restricting disk and network I/O, minimizing effect of security exploits
 - **Plug-in** process for each type of plug-in



Interprocess Communication

- Processes within a system may be **independent** or **cooperating**
- Cooperating process can affect or be affected by other processes, including sharing data
- Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
- Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Communications Models



Producer-Consumer Problem

- Paradigm for cooperating processes, *producer* process produces information that is consumed by a *consumer* process
- A buffer of items that can be filled by the producer and emptied by the consumer
 - *unbounded-buffer* places no practical limit on the size of the buffer
 - *bounded-buffer* assumes that there is a fixed buffer size

Bounded-Buffer Shared-Memory Solution

- Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

Bounded-Buffer – Producer Process

```
while (true) {  
    /* Produce an item */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing -- no free buffers */  
    buffer[in] = item;  
    in = (in + 1) % BUFFER_SIZE;  
}
```

Bounded Buffer – Consumer Process

```
while (true) {  
    while (in == out)  
        ; // do nothing-- nothing to consume  
  
    // remove an item from the buffer  
    item = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    return item;  
}
```

Question

- Which of following is true about the just discussed shared-memory solution for the producer-consumer problem with BUFFER_SIZE = 10 ?

A	The solution is incorrect because producer and consumer may modify the same entry at the same time
B	The solution is correct and maximum number of items in the buffer is 10.
C	The solution is correct but maximum number of items in the buffer is 9.
D	The solution is correct only if the producer runs faster than the consumer.
E	None of the above

Examples of IPC Systems – POSIX Shared Memory

- POSIX Shared Memory use memory-mapped files
 - Process first creates shared memory segment
`shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);`
 - Also used to open an existing segment to share it
 - Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Get pointer to the shared memory object

```
ptr = mmap(0, 4096, PROT_WRITE, MAP_SHARED, shm_fd, 0)
```

- Now the process could write to the shared memory

```
sprintf(ptr, "Writing to shared memory");
```

IPC POSIX Producer

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```

IPC POSIX Consumer

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Interprocess Communication – Shared Memory

- An area of memory shared among the processes that wish to communicate
- The communication is under the control of the users processes not the operating system.
- Major issues is to provide mechanism that will allow the user processes to synchronize their actions when they access shared memory.
- Synchronization is discussed in great details later.

Interprocess Communication – Message Passing

- Mechanism for processes to communicate and to synchronize their actions
- Message system – processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - **send**(*message*) – message size fixed or variable
 - **receive**(*message*)

Implementation Questions

- If P and Q wish to communicate, they need to:
 - establish a *communication link* between them
 - exchange messages via send/receive
- Implementation options:
 - How are links established?
 - Can a link be associated with more than two processes?
 - How many links can there be between every pair of communicating processes?
 - What is the capacity of a link?
 - Is the size of a message that the link can accommodate fixed or variable?
 - Is a link unidirectional or bi-directional?

Message Passing (Cont.)

- Implementation of communication link
 - Physical:
 - Shared memory
 - Hardware bus
 - Network
 - Logical:
 - Direct or indirect
 - Synchronous or asynchronous
 - Automatic or explicit buffering

Direct Communication

- Processes must name each other explicitly:
 - **send** (P , $message$) – send a message to process P
 - **receive**(Q , $message$) – receive a message from process Q
- Properties of direct communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

- Messages are directed and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox
- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication

- Operations
 - create a new mailbox (port)
 - send and receive messages through mailbox
 - destroy a mailbox
- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox A
 - receive**(*A, message*) – receive a message from mailbox A

Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 sends; P_2 and P_3 receive
 - Who gets the message?
- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.

Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null
- Different combinations of send() and receive() are possible
 - If both send and receive are blocking, we have a **rendezvous**

Producer – Consumer using blocking send & receive

Producer-consumer becomes trivial using blocking send() and receive()

Producer

```
while (true) {  
    item = Produce();  
    send(item);  
}
```

Consumer

```
while (true) {  
    receive(&item);  
  
    Consume(item);  
}
```

This code is correct and relatively simple.
Why don't we always just use message passing (vs shared memory)

```
/* W/O SHARED MEMORY */
```

Producer

```
int item;
```

```
while (TRUE) {  
    item = Produce ();  
    send (Consumer, &item);  
}
```

Consumer

```
int item;
```

```
while (TRUE) {  
    receive (Producer, &item);  
    Consume (item);  
}
```

- A. Message passing copies more data.
- B. Message passing only works across a network.
- C. Message passing is a security risk.
- D. We usually do use message passing!

Buffering

- Queue of messages attached to the link;
- Implemented in one of three ways
 1. Zero capacity – 0 messages are queued on a link.
Sender must wait for receiver (rendezvous).
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits

Unix Pipes

- Unix pipes allow communication between two processes using a buffer.


```
ls -l | grep test
```

- One process writes to the buffer and one process reads from the buffer.
- The communication is unidirectional.
- The buffer is accessed using a file descriptor.
- A pipe is created by using the pipe system call

```
int pipe(int* filedes);
```

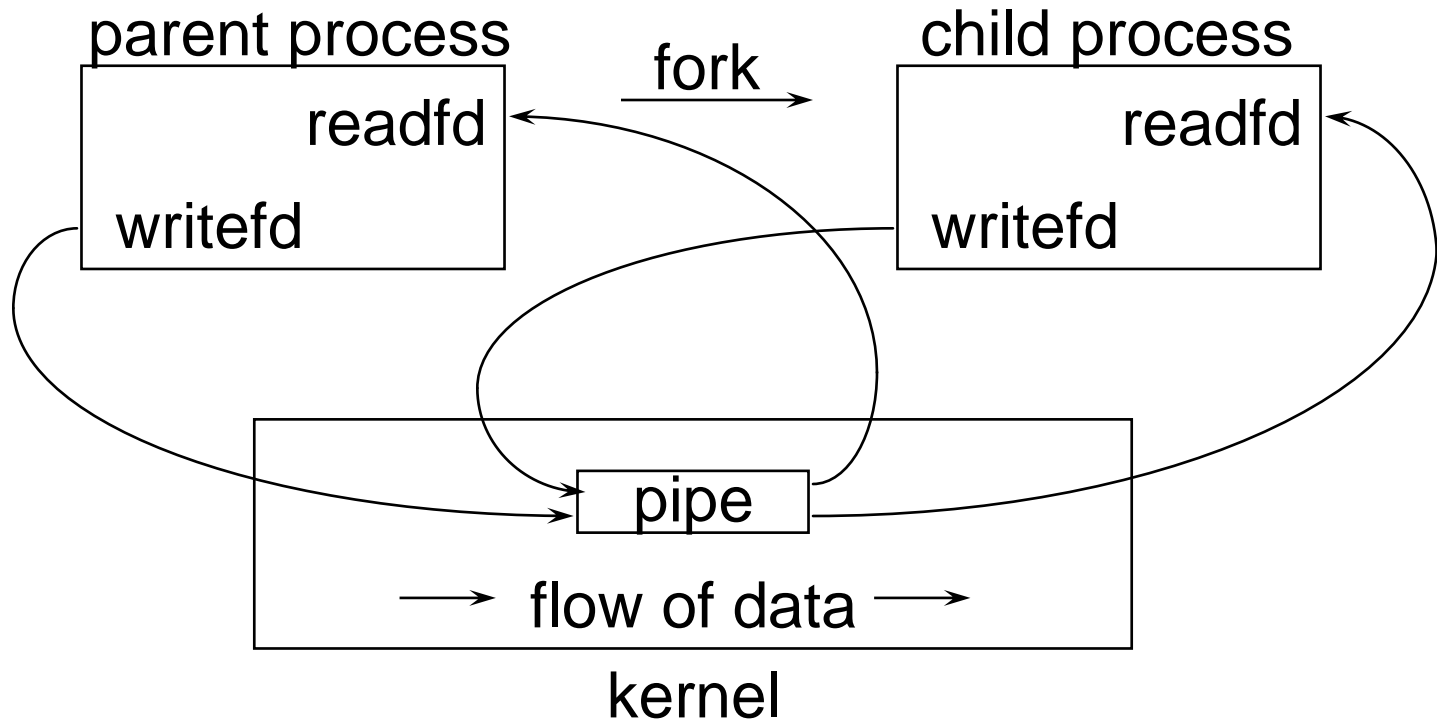
- Two file descriptors are returned
 - filedes[0] is open for reading
 - filedes[1] is open for writing
- The difference between a file and a pipe: pipe is a data structure in the kernel.
- Typical size is 512 bytes (Minimum limit defined by POSIX)

A Pipe Object



Ordinary Pipe

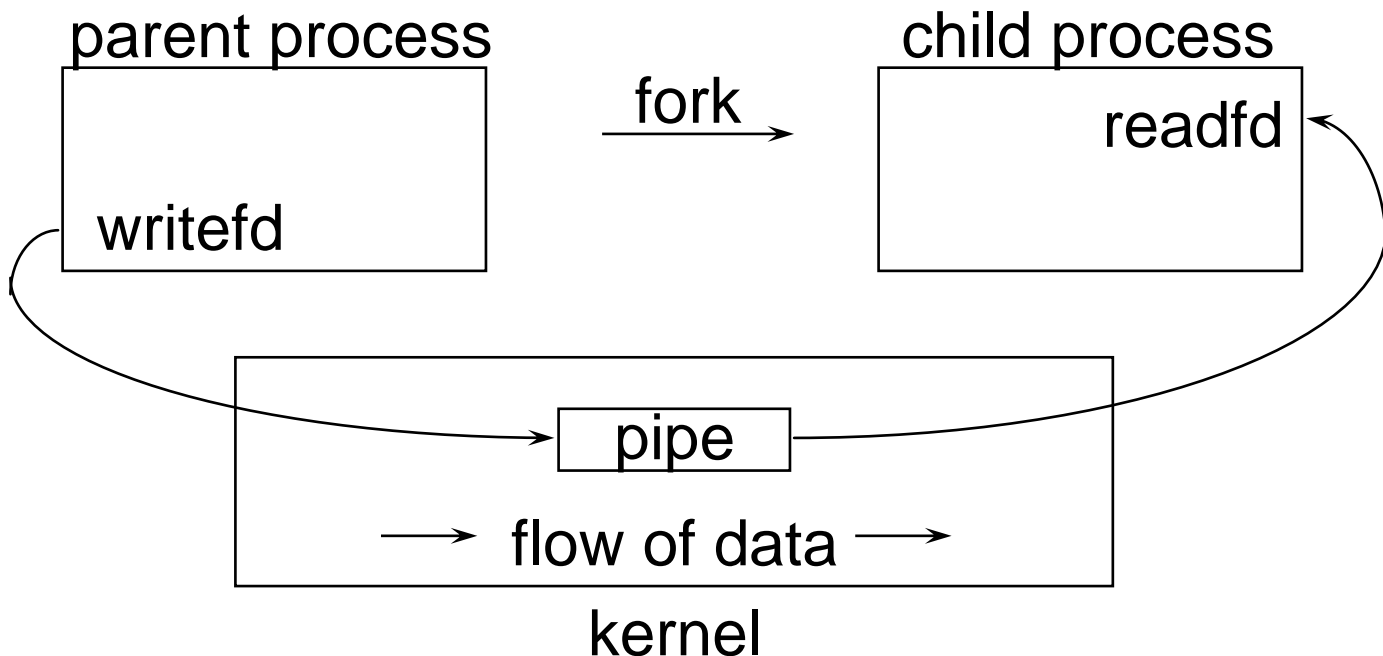
- First, a process creates a pipe, and then forks to create a copy of itself.



Pipe Examples

Parent writes file, child reads file

- parent closes read end of pipe
- child closes write end of pipe



Pipe Example

```
int main(void)
{
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    pid_t pid;
    int fd[2];

    /* create the pipe */
    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    printf("readfd = %d, writefd = %d\n", fd[0], fd[1]);

    /* now fork a child process */
    pid = fork();
    if (pid < 0) {
        fprintf(stderr, "Fork failed");
        return 1;
    }
}
```


Pipe Example (cont.)

```
if (pid > 0) { /* parent process */
    /* close the unused end of the pipe */
    close(fd[READ_END]);
    /* write to the pipe */
    write(fd[WRITE_END], write_msg, strlen(write_msg)+1);
    /* close the write end of the pipe */
    close(fd[WRITE_END]);
} else { /* child process */
    /* close the unused end of the pipe */
    close(fd[WRITE_END]);
    /* read from the pipe */
    read(fd[READ_END], read_msg, BUFFER_SIZE);
    printf("child read %s\n", read_msg);

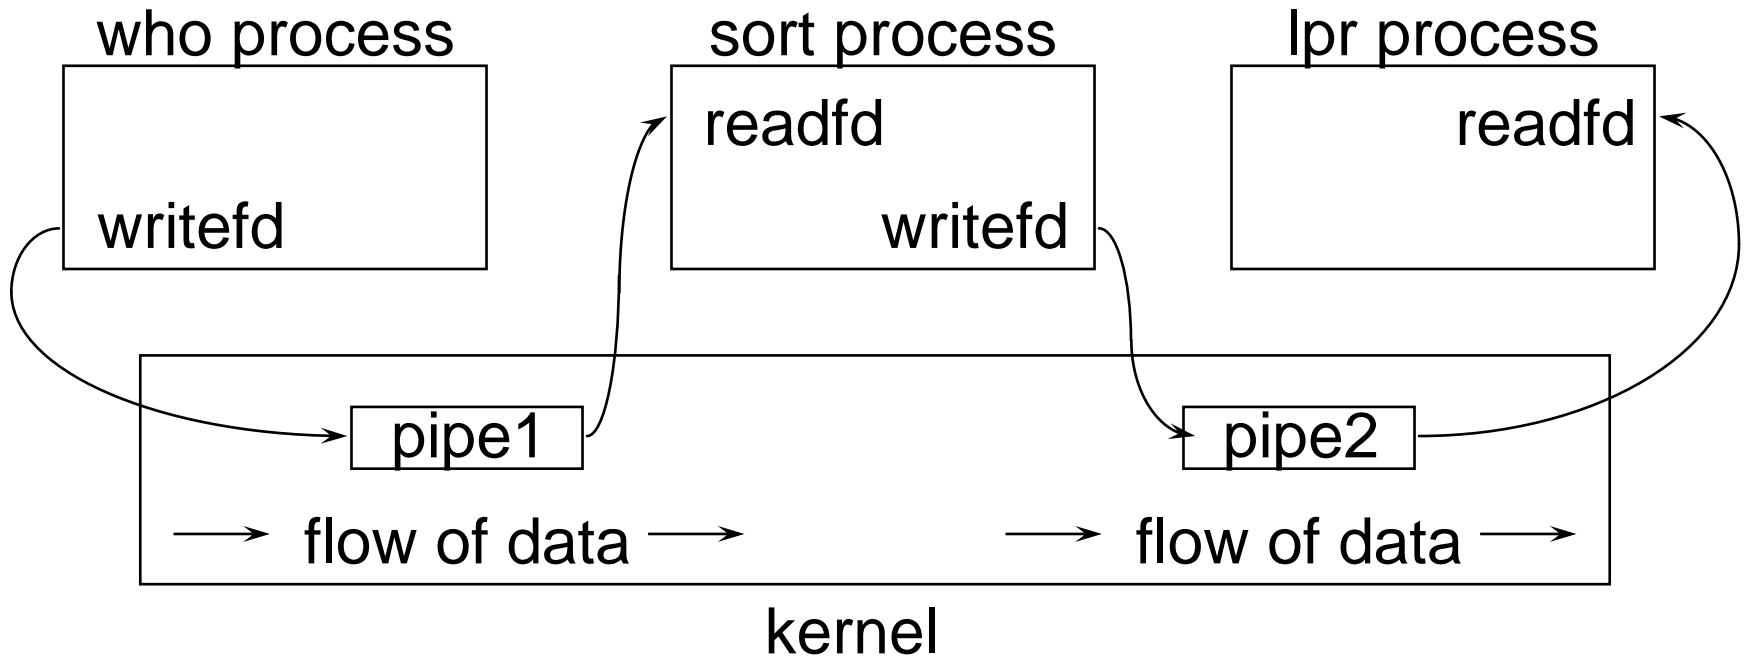
    /* close the read end of the pipe */
    close(fd[READ_END]);
}
```

If $fd[0] = 3$ and $fd[1] = 4$, what is the output of the above code?

Concatenated Pipe Examples

who | sort | lpr

- who process writes to pipe1
- sort process reads from pipe1, writes to pipe2
- lpr process reads from pipe2




Client-Server Communication

- Sockets
- Remote Procedure Calls

Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- All ports below 1024 are ***well known***, used for standard services
- Communication consists between a pair of sockets



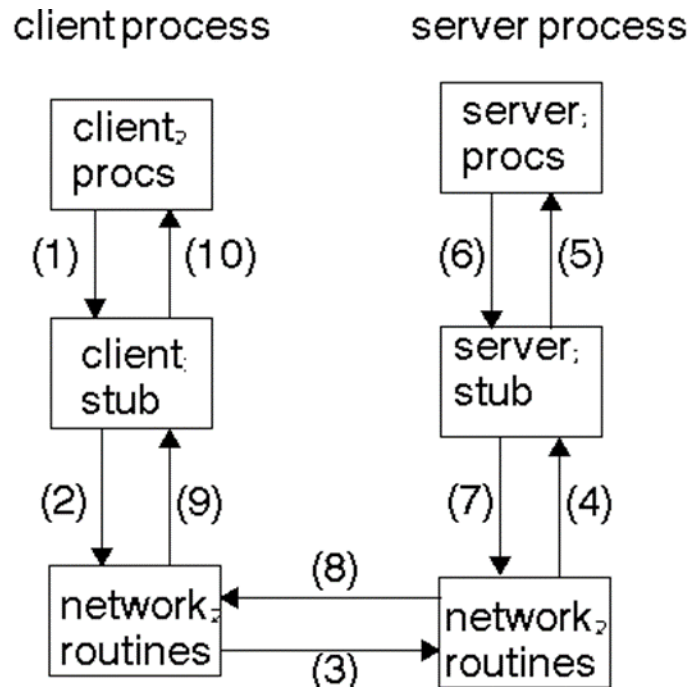
Date Server

```
public static void main(String[] args) {  
    try {  
        ServerSocket sock = new ServerSocket(6013);  
        // now listen for connections  
        while (true) {  
            Socket client = sock.accept();  
            // we have a connection  
            PrintWriter pout = new PrintWriter(client.getOutputStream(),  
true);  
            // write the Date to the socket  
            pout.println(new java.util.Date().toString());  
            // close socket and resume listening for more connections  
            client.close();  
        }  
    } ...  
}
```

Date Client

```
public static void main(String[] args)  {
    try {
        // this could be changed to an IP name or address
        // other than the localhost
        Socket sock = new Socket("127.0.0.1",6013);
        InputStream in = sock.getInputStream();
        BufferedReader bin = new BufferedReader(new
        InputStreamReader(in));
        // read data from the socket
        String line;
        while( (line = bin.readLine()) != null)
            System.out.println(line);
        sock.close();
    }
}
```

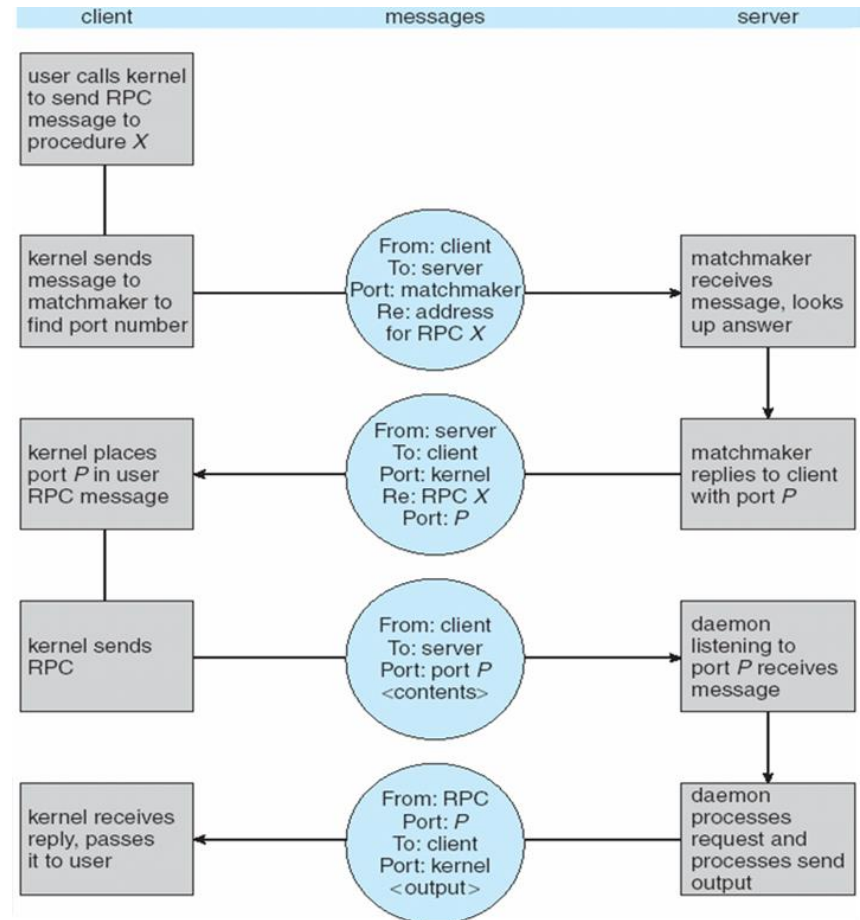
Remote Procedure Calls



- Remote procedure call (RPC) abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshalls* the parameters.
 - Data representation handled via **External Data Representation (XDR)** format
- The server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.
- On Windows, stub code compile from specification written in **Microsoft Interface Definition Language (MIDL)**

Execution of RPC

- Remote communication has more failure scenarios than local
 - Messages can be delivered ***exactly once*** rather than ***at most once***
- OS typically provides a rendezvous (or **matchmaker**) service to connect client and server



Summary

- Process is a program in execution that can be in a number of states
 - New, running, waiting, ready, terminated
- Process creation and termination
- Inter-process communications
 - Shared memory, and message passing
- Client-server communication
 - Socket, RPC, ...